

# EOOP 20L Preliminary project - Neural Network

Marcin Wojnarowski (303880)

*Disclaimer:* This document is not meant to be an explanation of the inner workings of a Neural Network. It will strictly focus on the interface and implementation.

## 1. Description

This project will focus on implementing a basic multilayer perceptron (known as an ANN: Artificial Neural Network) stripped of all of the more advanced concepts like normalizing weights, mini-batch training, or vectorized computation. The implementation will consist of 5 classes in total, each handling a separate part of what a Neural Network forms.

As a Proof of Concept I will train a NN to solve the [XOR](#) problem: it is known as a simple non-linear problem that requires at least one hidden layer in order to be solved, thus solving it proves that a given NN is able to solve other arbitrary non-linear problems. More or less effectively, but still. If time allows, the NN will be also put against the [MNIST database](#).

### Classes

**Serializer** This class only exists as a base of a different one; it is an abstract class. It requires that deriving classes have implemented a `serialize` method and a static `deserialize` method. Then, the **Serializer** is able to provide such methods as `from_file` or `to_file`, allowing for easy serialization handling no matter who is the parent class.

This class will prove to be very useful when saving the NN's weights to a file and then loading them back in whenever needed. This allows for an interruptible training cycle. Additionally it unlocks the ability of configuring the NN from a text file rather than fiddling with code.

**Config** **Config** conforms **Serializable**<sup>1</sup>. Stores the learning rate and the sizes of the network: input, output, hidden layers, hidden neurons.

---

<sup>1</sup>Serializable: a class that implements the `serialize` and `deserialize` methods

**Matrix** The main purpose of the `Matrix` class is to simplify matrix operations. It will override the 4 basic arithmetic operators `*/+` providing a layer of abstraction when using 2d arrays. `*/+` will work for both scalars and other matrices while `/` will be constrained to just scalars. Additionally there will be 2 other linear algebraic operations: transposition and the dot product. Cross product will be omitted on purpose: while it would be great for completeness, the usefulness is minimal.

`Matrix` conforms `Serializable` as well.

**NNFunctions** `NNFunctions` stores function used by the NN: activation function, derivative of the activation function, output layer function, derivative of the output layer function, and cost function. Constructing a `NNFunction` object consists of providing the previously mentioned functions, or choosing ready functions from an enum.

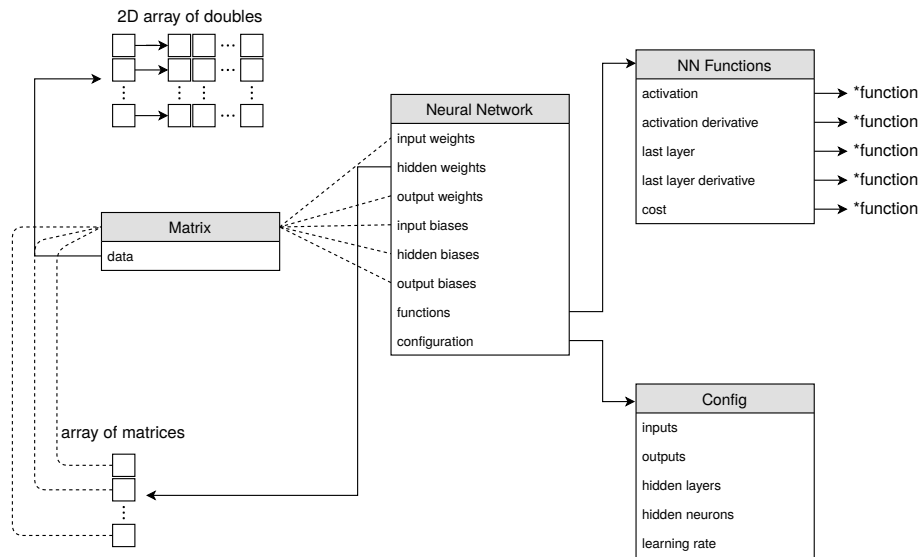
`NNFunctions` is `Serializable` only if the chosen functions are from the preset, custom ones can't be serialized.

**NeuralNetwork** This is the *brain* class. Uses all of the classes above to construct a friendly interface for training and performing guesses.

## 2. Memory map

The neural network owns all of its data (except pointers to functions which have static lifetimes).

Dotted lines indicate that the data has a structure of the connected type (they don't all point to the same `Matrix` instance).



### 3. Class declarations

The lack of raw pointers is a conscious decision. They are a source of bugs and if a need for pointers will arise, smart pointers will be used instead. Because all properties are auto-cleaned/copied no destructors or copy constructors are present.

#### Serializer

```

template <typename T>
class Serializer {
protected:
    /// protected on purpose, Serializer is an abstract class
    Serializer();

public:
    /// deserializes a file into the parent object.
    /// throws if file does not exist
    static auto from_file(const string& path) -> T;

    /// takes a path and serializes the parent into the pointed file.
    /// overwrites all content if the file already exists
    auto to_file(const string& path) const -> void;

```

```

    /// prints the serialized object
    // friend auto operator<<<>(ostream& os, const T& obj) -> ostream&;

    /// virtual methods that have to be implemented by parent classes
    /// then the serializer can work properly
    virtual auto serialize() const -> string = 0;
    /// this is impossible in c++, therefore it only serves purely as
    /// documentation
    // virtual static auto deserialize(const string& str) -> T = 0;
};

```

## Config

```

class Config : Serializer<Config> {
public:
    /// properties of a neural network
    /// because they are constant, there is no need encapsulating them
    const unsigned int inputs, outputs, layers, hidden_neurons;
    const double learning_rate;

    /// constructor accepting all 4 parameters
    Config(const unsigned int inputs, const unsigned int outputs,
           const unsigned int layers, const unsigned int hidden_neurons,
           const double learning_rate);

    /// overriding the virtual methods of Serializer
    auto serialize() const -> string override;
    static auto deserialize(const string& str) -> Config;
};

```

## Matrix

```

class Matrix : Serializer<Matrix> {
public:
    /// size of the matrix
    const unsigned int rows, columns;

    /// constructor takes the dimensions of the matrix
    Matrix(const unsigned int rows, const unsigned int columns);

    /// randomizes the matrix with a given range
    auto randomize(const double min = -1.0, const double max = 1.0) -> void;

    /// operator overloads for matrix operations

```

```

    /// if on both sides of the operation theres a matrix then the operation is
    /// done element wise, unless it is * where a matrix multiplication is
    /// performed instead
    /// in-place
    auto operator+=(const Matrix& rhs) -> Matrix&;
    auto operator+=(const double& rhs) -> Matrix&;
    auto operator-=(const Matrix& rhs) -> Matrix&;
    auto operator-=(const double& rhs) -> Matrix&;
    auto operator*=(const double& rhs) -> Matrix&;
    auto operator/=(const double& rhs) -> Matrix&;
    /// global
    friend auto operator+(const Matrix& lhs, const Matrix& rhs) -> Matrix;
    friend auto operator+(const Matrix& lhs, const double& rhs) -> Matrix;
    friend auto operator-(const Matrix& lhs, const Matrix& rhs) -> Matrix;
    friend auto operator-(const Matrix& lhs, const double& rhs) -> Matrix;
    friend auto operator*(const Matrix& lhs, const Matrix& rhs) -> Matrix;
    friend auto operator*(const Matrix& lhs, const double& rhs) -> Matrix;
    friend auto operator/(const Matrix& lhs, const double& rhs) -> Matrix;
    /// indexing
    auto operator[](size_t idx) const -> vector<double>&;

    /// transposing flips the x and y axis
    auto transpose() const -> Matrix;

    /// overriding the virtual methods of Serializer
    auto serialize() const -> string override;
    static auto deserialize(const string& str) -> Matrix;

private:
    /// thats where the data is stored. Vector was chosen because while the
    /// size is immutable and array would seem like a more fitting choice, vector
    /// provides a much safer interface with negligible overhead
    vector<vector<double>> data_;
};

```

## NNFunctions

```

class NNFunctions : Serializer<NNFunctions> {
public:
    /// enums listing available function
    /// __custom means the function was provided
    enum class Activation { sigmoid, relu, tanh, __custom };
    enum class LastLayer { softmax, __custom };
    enum class Cost { mean_square, __custom };

```

```

    /// type definitions of the functions
    /// a function that takes a double and decides if its active
    typedef auto (*Activating)(double) -> double;
    /// a function that takes an array of doubles and maps it to different values
    typedef auto (*Mapping)(vector<double>) -> vector<double>;
    /// a function that takes an array of doubles and reduces it to a single value
    typedef auto (*Reducing)(const vector<double>&) -> double;

    /// collection of functions
    const Activating activation, d_activation;
    const Mapping last_layer, d_last_layer;
    const Reducing cost;

    /// constructor accepting enums describing pre-made functions
    NNFunctions(Activation af, LastLayer llf, Cost cf);
    /// constructor accepting functions
    NNFunctions(const Activating af, const Activating daf, const Mapping llf,
               const Mapping dllf, const Reducing cf);

    /// overriding the virtual methods of Serializer
    auto serialize() const -> string override;
    static auto deserialize(const string& str) -> NNFunctions;

private:
    /// remembering which functions were chosen, this information is needed for
    /// serialization
    Activation af_;
    LastLayer llf_;
    Cost cf_;
};

```

## NeuralNetwork

```

class NeuralNetwork : Serializer<Matrix> {
public:
    /// constructor takes the previously defined configuration
    NeuralNetwork(Config config, NNFunctions funcs);

    /// performs a classification guess, it is not meant for regression problems
    auto guess(const Matrix& inputs) const -> unsigned int;

    /// trains the network `n` amount of times using online training
    /// inputs and expected have to be linearly aligned: first element of inputs
    /// have to correspond to first element from expected and so on
    auto train(const vector<Matrix>& inputs, const vector<Matrix>& expected,

```

```

        unsigned int n) -> void;

    /// overriding the virtual methods of Serializer
    auto serialize() const -> string override;
    static auto deserialize(const string& str) -> NeuralNetwork;

private:
    /// weights of the connections
    Matrix input_w_;
    vector<Matrix> hidden_w_;
    Matrix output_w_;

    /// biases of the neurons
    Matrix input_b_;
    Matrix hidden_b_;
    Matrix output_b_;

    /// functions
    NNFunctions funcs_;

    /// config
    Config config_;

    /// sends inputs through the whole network and returns the output layer
    auto feedforward(const Matrix& inputs) const -> Matrix;

    /// backpropagates the expected output from some input, adjusts the weights,
    /// then returns the cost of the network
    auto backpropagate(const Matrix& inputs, const Matrix& expected) -> double;
};

```

## 4. Demos

Objects are easily saveable

```

Matrix m(10, 12);
m.save_to("/path/to/file");

```

Not all of them though

```

NNFunctions f(NNFunctions::Activation::relu, NNFunctions::LastLayer::__custom,
              NNFunctions::Cost::mean_square);

```

```

m.serialize(); // throws an exception, a custom function was used      ^^^^^^^^^

```

Matrices can be multiplied only if the sizes align

```
Matrix m1(1, 2);
Matrix m2(2, 2);
Matrix m3(3, 2);
```

```
m1 * m2; // ok
m1 * m3; // error!
```

They can be also scaled/moved by scalars

```
Matrix m(10, 10);
```

```
m += 1;
m -= 2;
m *= 3;
m /= 4;
```

Or added element wise

```
Matrix m1(2, 2);
Matrix m2(2, 2);
Matrix m3(3, 2);
```

```
m1 += m2;
m1 -= m2;
m1 += m3; // error, sizes do not align
```

Elements can be accessed (for reading and writing) with the index operator

```
Matrix m(3, 2);
```

```
m[1][0] = 1.2;
```

```
Matrix m2 = m.transpose();
```

```
assert(m2[0][1] == 1.2); // ok
```

The neural network provides a very high level API, only `guess` and `train` methods are available

```
// XOR example
auto nn = NeuralNetwork::from_file("./backup");
```

```
vector<Matrix> inputs;
vector<Matrix> expected;
```

```
// false false -> false
{
    Matrix i(2, 1);
    i[0][0] = 0.0;
    i[1][0] = 0.0;
```



```

        inputs.push_back(i);

        Matrix e(2, 1);
        e[0][0] = 1.0;
        e[1][0] = 0.0;
        expected.push_back(e);
    }

    // false true -> true
    {
        Matrix i(2, 1);
        i[0][0] = 0.0;
        i[1][0] = 1.0;
        inputs.push_back(i);

        Matrix e(2, 1);
        e[0][0] = 0.0;
        e[1][0] = 1.0;
        expected.push_back(e);
    }

    // true false -> true
    {
        Matrix i(2, 1);
        i[0][0] = 2.0;
        i[1][0] = 0.0;
        inputs.push_back(i);

        Matrix e(2, 1);
        e[0][0] = 0.0;
        e[1][0] = 1.0;
        expected.push_back(e);
    }

    // true true -> false
    {
        Matrix i(2, 1);
        i[0][0] = 1.0;
        i[1][0] = 1.0;
        inputs.push_back(i);

        Matrix e(2, 1);
        e[0][0] = 1.0;
        e[1][0] = 0.0;
        expected.push_back(e);
    }
}

```

```
nn.train(inputs, expected, 1000);
```

```
assert(nn.guess(inputs[0]) == 0);  
assert(nn.guess(inputs[1]) == 1);  
assert(nn.guess(inputs[2]) == 1);  
assert(nn.guess(inputs[3]) == 0);
```

There are unit tests for every method. They can be found [on GitHub](#).

- 
- The code is formatted using `clang-format` with the `Google` preset
  - Compiled with `g++` version 9.2 with the `-std=c++2a` flag
  - Naming convention:
    - Type aliases, classes, structs, enums, concepts: `PascalCase`
    - Private fields: `snake_case` with an underscore at the end
    - All of the rest: `snake_case`
  - Return types will be annotated with the `auto <name>() → <type>` syntax
  - Project is hosted [on GitHub](#)